

# Big Data Project - Yelp Dataset


## Data Aggregation and Recommendation System Implementation

Truc Cao

Tejinder Kaur

Amman Shareef

Kathia Teran



**Yelp Open Dataset**  
An all-purpose dataset for learning

The Yelp dataset is a subset of our businesses, reviews, and user data for use in personal, educational, and academic purposes. Available as JSON files, use it to teach students about databases, to learn NLP, or for sample production data while you learn how to make mobile apps.

**The Dataset**

Reviews	Businesses	Pictures	Metropolitan Areas
8,021,122 reviews	209,393 businesses	200,000 pictures	10 metropolitan areas

1,320,761 tips by 1,968,703 users  
Over 1.4 million business attributes like hours, parking, availability, and ambience  
Aggregated check-ins over time for each of the 209,393 businesses

# Problem and Motivation

The Restaurant market in New York is worth \$17 billion dollars, with currently 31,061 businesses. Having a dataset like that of Yelp's readily available online is of much value to any business trying to enter this market. There is a myriad of ways a business can benefit from a dataset like this, since it has detailed information on what customers like and dislike (among other information) about the way in which a restaurant business operates.

Yelp dataset has a lot of data but hard to extract useful information

We want to:

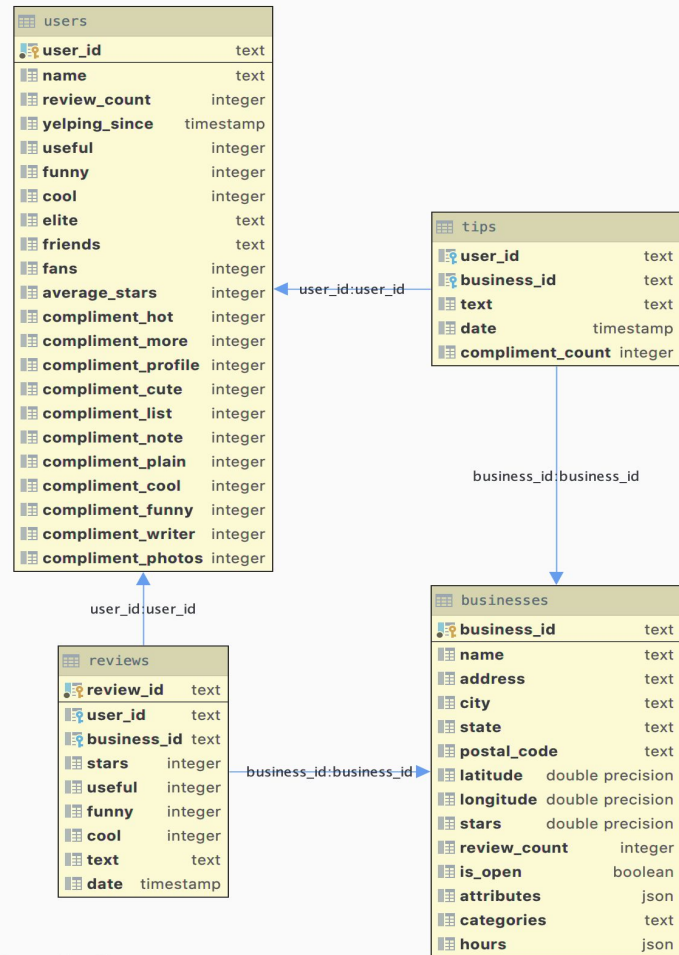
- Get statistics on restaurants (cuisines, popularity, ...)
- Recommend restaurants to customers to increase revenue
- Recommend users to restaurants so they can target them with advertising

Using:

- Spark (Big Data Framework) to cleanup, filter and aggregate data
- Research papers about recommendation system
- Spark ML (Machine Learning) to provide useful recommendation

# Yelp Dataset

- Yelp is a popular online crowd-sourced local business review platform.
- Sample: <https://www.yelp.com/dataset>
- The dataset contains:
  - 192,609 businesses
  - 1,637,138 users
  - 6,685,900 reviews
  - 1,223,094 tips
  - 161,950 check-ins
  - 36 states
  - 1,307 cities
  - All the data files are in JSON format.



# Samples

```
{
  "business_id": "0W27hbZN7Z-PrkhAb0y9Eg",
  "name": "B Montréal",
  "address": "1207-A Rue Rachel E",
  "city": "Montréal",
  "state": "QC",
  "postal_code": "H2J 2J8",
  "latitude": 45.5266477836,
  "longitude": -73.5735161233,
  "stars": 4.5,
  "review_count": 3,
  "is_open": 1,
  "attributes": {
    "RestaurantsPriceRange2": "1",
    "RestaurantsTableService": "False",
    "RestaurantsTakeOut": "True",
    "OutdoorSeating": "True",
    "RestaurantsReservations": "False",
    "Alcohol": "u'none'",
    "BikeParking": "True",
    "WheelchairAccessible": "True",
    "Caters": "True",
    "HasTV": "False",
    "GoodForMeal": "{ 'dessert': False, 'latenight': False, 'lunch': False, 'dinner': False, 'brunch': False, 'breakfast': False }",
    "GoodForKids": "True",
    "RestaurantsDelivery": "True",
    "WiFi": "u'free'"
  },
  "categories": "Coffee & Tea, Food, Juice Bars & Smoothies, Delis, Restaurants, Sandwiches",
  "hours": {
    "Wednesday": "9:0-20:0",
    "Thursday": "9:0-20:0",
    "Friday": "9:0-20:0",
    "Saturday": "8:0-16:0"
  }
}
```

business.json

```
{
  "business_id": "-Lw8Ve0NLbR0djHGw2fMOA",
  "date": "2014-06-30 22:57:27, 2014-06-30 22:58:28, 2017-05-19 18:24:18"
}
```

checkin.json

```
{
  "user_id": "zhseuNa_3b246gzcy8BXA",
  "name": "Briana",
  "review_count": 44,
  "yelping_since": "2009-10-16 23:54:29",
  "useful": 40,
  "funny": 3,
  "cool": 12,
  "elite": "",
  "friends": "_SEBcjOwne0V1VV_vESFQ_vtNhztatfsxCCsbC_JUK8Cw, EPyRgySYsr365gAgF2r4Ww, 3mNk60ynkQYRNJGf3YAq1A, NFU0zDaTMEQ4-X9dbQWd9A...X",
  "fans": 2,
  "average_stars": 4.07,
  "compliment_hot": 0,
  "compliment_more": 1,
  "compliment_profile": 0,
  "compliment_cute": 0,
  "compliment_list": 0,
  "compliment_note": 0,
  "compliment_plain": 0,
  "compliment_cool": 0,
  "compliment_funny": 0,
  "compliment_writer": 0,
  "compliment_photos": 0
}
```

user.json

```
{
  "review_id": "BTDBNxb7m6wuSTy09_Zz4A",
  "user_id": "oV4PUFp402brd3bGhu5cjg",
  "business_id": "m97jaBYRscg-hqDjMVIIng",
  "stars": 4,
  "useful": 0,
  "funny": 0,
  "cool": 0,
  "text": "This is our go-to place for lunch and just a friendly atmosphere. Very consistent food. A",
  "date": "2017-03-03 22:35:42"
}
```

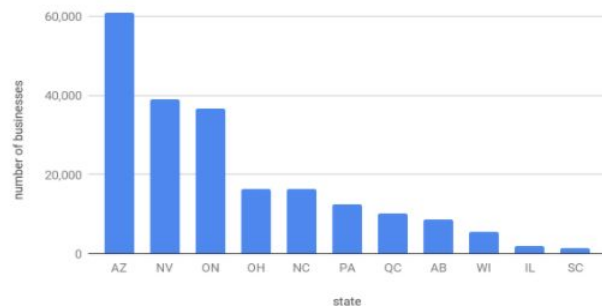
review.json

```
{
  "user_id": "ky3DB9i9lDJ70AZdkZyv7g",
  "business_id": "m9ybLDUrbagso1IT06bBLA",
  "text": "Excellent price on tires here!",
  "date": "2016-01-07 18:01:31",
  "compliment_count": 0
}
```


tip.json

# Dataset: Distribution


Number of businesses by state



Number of businesses by city



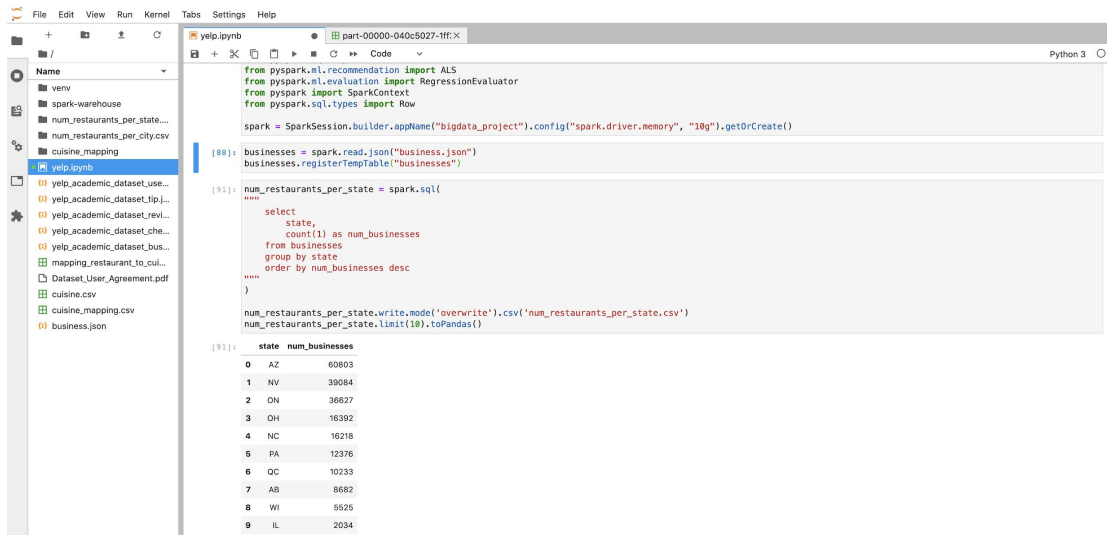
Number of Businesses by Category



- Out of 36 states (US/Canada), 11 has more than 1,000 businesses
- Out of 1,307 cities, 28 has more than 1000 businesses
- Restaurant is by far the most represented category

# Preparing the data

- Use Jupyter notebook and Spark SQL
- Jupyter notebook:
  - Interactive notebook in the browser
  - Can share easily with other people
- SparkSQL:
  - SQL queries are easier to write (and to read) than dataframe operations
  - Take advantage of Spark to scale to multiple machines to speed up the process



The screenshot shows a Jupyter notebook with a file explorer on the left and a code editor on the right. The file explorer lists various files, including 'yelp.ipynb'. The code editor contains the following Python code:

```
from pyspark.ml.recommendation import ALS
from pyspark.ml.evaluation import RegressionEvaluator
from pyspark import SparkContext
from pyspark.sql.types import Row

spark = SparkSession.builder.appName("bigdata_project").config("spark.driver.memory", "10g").getOrCreate()

[88]: businesses = spark.read.json("business.json")
      businesses.registerTempTable("businesses")

[91]: num_restaurants_per_state = spark.sql(
      """
      select
        state,
        count(1) as num_businesses
      from businesses
      group by state
      order by num_businesses desc
      """)

num_restaurants_per_state.write.mode('overwrite').csv('num_restaurants_per_state.csv')
num_restaurants_per_state.limit(10).toPandas()
```

The output of the last cell is a table showing the number of businesses per state:

	state	num_businesses
0	AZ	60803
1	NV	39084
2	ON	36627
3	OH	16392
4	NC	16218
5	PA	12376
6	QC	10233
7	AB	8682
8	WI	5525
9	IL	2034

# Preparing the data

## Steps:

- Only keep restaurants: “Restaurants”
- Break down categories (one row per category):
  - “Ramen, Sushi” becomes 2 rows for each
- Recategorize using 32 kind of cuisines (American, Korean):
  - Use manually created mapping (114 mappings)
  - In case of multiple cuisine we choose one arbitrarily
- Only keep restaurants with a cuisine
- Only keep cities that have over 200 restaurants
- Only keep reviews and users associated to them

	cuisine	category
0	American	American (Traditional)
1	American	American (New)
2	American	Steakhouses
3	American	Bagels
4	American	Cajun/Creole
...	...	...
109	Asian	Cambodian
110	Korean	Korean
111	Asian	Laotian
112	Asian	BurmesePizza
113	Thai	Thai

cuisine.csv (Cuisine mapping)

## Initialization

```
spark = SparkSession.builder.appName("yelp").config("spark.driver.memory", "10g").getOrCreate()
businesses = spark.read.json("yelp_academic_dataset_business.json")
businesses.registerTempTable("businesses")
```

	address	attributes	business_id	categories	city	hours	is_open	latitude	longitude	name
0	404 E Green St	(None, None, 'none', None, None, None, None, F...	pQeaRpvuhoEqudo3uymHIQ	Ethnic Food, Food Trucks, Specialty Food, Impo...	Champaign	(11:30-14:30, 11:30-14:30, None, None, 11:30-1...	1	40.110446	-88.233073	The Empanadas House
1	4508 E Independence Blvd	(None, None, None, None, None, None, None, Non...	CsLQLiRoafpJPJSKNX2h5Q	Food, Restaurants, Grocery, Middle Eastern	Charlotte	None	0	35.194894	-80.767442	Middle East Deli
2	15480 Bayview Avenue, unit D0110	(None, None, u'none', None, None, None, None, ...	eBEfgOPG7pvFhb2wcG9I7w	Restaurants, Cheesesteaks, Poulineries	Aurora	(11:0-22:0, 11:0-22:0, 11:0-22:0, 11:0-21:0, 1...	1	44.010962	-79.448677	Philly Philly's

## Break down categories:

```
restaurant_categories_df = spark.sql("""
SELECT
    business_id,
    explode(split(categories, ',')) as category
FROM businesses
where categories like '%Restaurants%'
""")
restaurant_categories_df.registerTempTable('restaurant_categories')
```

	business_id	category
0	AtD6B83S4MbmQ0t7iDnUVA	Sushi Bars
1	AtD6B83S4MbmQ0t7iDnUVA	Dim Sum
2	AtD6B83S4MbmQ0t7iDnUVA	Ramen

## Map to cuisines:

```
cuisine_mapping_df = spark.read.csv("cuisine.csv", header=True)
cuisine_mapping_df.registerTempTable("cuisine_mapping")

restaurant_cuisines_df = spark.sql("""
SELECT
    business_id,
    LAST(cuisine) as cuisine
FROM restaurant_categories
JOIN cuisine_mapping USING (category)
GROUP BY business_id
""")
restaurant_cuisines_df.registerTempTable('restaurant_cuisines')
```

	business_id	cuisine
	--9e10NYQuAa-CB_Rrw7Tw	American
	-VASjhmAbKF3Pb_-8rh3xg	Coffee & Tea
	-cxD1NimFldATDUsN-oa3A	Mexican
	-r8SVtXXG6_T3mP5GXRAW	Coffee & Tea
	0859wfd1BQH4G46Zpwhc0ZQ	Bars
	09OYbFNrS1n8u5gE6W9ItA	German
	0DwMrcy7_X_C_mP8_QcXug	American
	0bqV9uzFVz98Bn_RlmcJTg	Mexican
	0owIRP_z5RcYKKmh5RnD7A	Fast Food
	1NmGVWYIF4iMngM6arKJTQ	Vietnamese

# Final dataset

- 192,609 businesses
- 1,637,138 users
- 6,685,900 reviews
- 1,223,094 tips
- 161,950 check-ins
- 36 states
- 1,307 cities




- 41,029 restaurants (-78%)
- 1,006,767 users (-39%)
- 3,447,322 reviews (-48%)
- 668,619 tips (-45%)
- 39,896 check-ins (-75%)
- 10 states -72%()
- 30 cities (-98%)





# Computing different aggregations of the data



- Number of restaurants that are still active/closed for each dimension
- Number of restaurants /ratings / city
- Most popular restaurant and influencer for dimension (rating, cuisine, ...)
- Do all the aggregations using Spark SQL




Restaurant by Status



Restaurant by Status




Restaurant by Status




• All,

Restaurant by Rating




las vegas,

Restaurant by Rating



toronto

Restaurant by Rating



# Timeseries data

## Number of checkins / cuisine / city over time

### Map tables from file

```
restaurants_df = spark.read.json("final_restaurants_all.json")
restaurants_df.registerTempTable('restaurants')

checkins_raw_df = spark.read.csv("final_restaurant_checkin_all.json")
checkins_raw_df.registerTempTable('checkins_raw')
```

### Master aggregation by city/state, cuisine, month


```
checkins_by_date_df = spark.sql("""
WITH checkin_date AS (
  SELECT
    business_id,
    explode(split(date, ',')) as date
  FROM checkins
)
SELECT
  city,
  state,
  cuisine,
  substring(date, 0, 7) as date,
  count(1) as num_checkins
FROM checkin_date
JOIN restaurants USING (business_id)
GROUP BY city, state, cuisine, date
ORDER BY date
""")
```

### Pivot table for number of checkins for each cuisine in Toronto over time

```
popular_cuisines = ['American', 'Bars', 'Japanese', 'Chinese', 'Mexican', 'Fast Food', 'Coffee & Tea', 'Italian', 'Pizza', 'Mediterranean', 'Thai',
                    'Vietnamese', 'Korean', 'French', 'Indian']
tmp_checkins_cuisine_toronto_df = checkins_by_date_df.filter(f.col('city') == 'Toronto').filter(f.col('cuisine').isin(popular_cuisines)) \
    .groupBy('date').pivot('cuisine').sum('num_checkins').sort('date')
checkins_cuisine_toronto_df = tmp_checkins_cuisine_toronto_df.na.fill(0)
checkins_cuisine_toronto_df.write.mode('overwrite').csv('checkins_cuisine_toronto.csv')
```

```
checkins_cuisine_toronto_df.toPandas()
```

	date	American	Bars	Chinese	Coffee & Tea	Fast Food	French	Indian	Italian	Japanese	Korean	Mediterranean	Mexican	Pizza	Thai	Vietnamese
0	2010-01	86	37	8	22	3	4	6	19	8	2	4	14	5	9	4
1	2010-02	132	50	27	41	4	8	19	32	35	6	5	18	6	17	10
2	2010-03	212	88	32	81	21	10	14	39	45	10	6	40	3	25	18
3	2010-04	215	144	35	80	8	20	16	45	47	11	9	36	7	22	13
4	2010-05	245	113	30	79	18	17	16	60	37	9	12	35	13	27	10



Number of checkins by cuisine over time in Toronto

# Influencers: how top reviewers rate businesses

- Influencers being those users with the highest amount of: friends, fans, useful rating, elite status, and oldest accounts ('yelping since').

```
# Create a SparkSession (the entry point to using the database)
spark = SparkSession.builder.appName("Popular_Reviewer").getOrCreate()

# Get the raw data CONVERTS JSON, used to be 'lines'
business_dataFrame = spark.read.json(path1)
reviews_dataFrame = spark.read.json(path2)
users_dataFrame = spark.read.json(path3)

# Some SQL-style magic to sort all movies by popularity in one line!
top_business = business_dataFrame.select("business_id", "name", "city").orderBy("city", ascending=False)
top_reviews = reviews_dataFrame.select("business_id", "user_id").orderBy("user_id", ascending=False)
top_reviewer = users_dataFrame.select("user_id", "name", "review_count").orderBy("review_count", ascending=False)

top_reviewers = top_business.join(top_reviews, on=['business_id'], how='inner').orderBy("city", ascending=False).\
    join(top_reviewer, on=['user_id'], how='inner').orderBy("review_count", ascending=False)

top_reviewers.show(10, False)
```

# Most Popular Influencers

## I. BASED ON CITY (considering Las Vegas as the most yelped city for restaurants)

- Restaurant and city with the highest review count:

business_id	name	city	review_count
4JNXUY8wbaaDmk38Pz1w	Mon Ami Gabi	Las Vegas	8348
f4x1YBxkLrZg652xt2KR5g	Hash House A Go Go	Las Vegas	5763
cYwJAZA6I12XNkm2rtXd5g	Gordon Ramsay BurGR	Las Vegas	5484

### TOP

- 'Useful' influencer in Las Vegas (review count is from the business)

user_id	business_id	name	city	review_count	name useful
--2vR0DIsmQ6Wfc5zKWigw	IB8z1lGra0g9LU7qVLPyg	Fashion Show	Las Vegas	739	Harald 154202

- 'Oldest (yelping since)' influencer in Las Vegas (business review count)

user_id	business_id	name	city	review_count	name yelping_since
lnkN_dof3J9xekchVC-v68A	ubz4CaZagQuGv2N9gFAdw	Botero	Las Vegas	434	Jeremy 2004-10-12 08:46:43

- Influencer with the most 'Fans' in Las Vegas (business review count)

user_id	business_id	name	city	review_count	name fans
37cpUoM8h1kS0fReIEBd-Q	0qet57CmMA5qUm6gPFUTpg	Di Fara Pizza	Las Vegas	150	Mike 9538

- Influencer with the most 'review counts' in Las Vegas

user_id	business_id	name	city	name review_count
8k3a0-mPeyhbRSHUuca5aA	607-wkCpc1KF75JZLOTcMw	Circus Circus Las Vegas Hotel and Casino	Las Vegas	Victor 13278

## II. BASED ON CATEGORY:

- Influencer with the most 'Fans'

user_id	business_id	name	city	categories	name fans
37cpUoM8h1kS0fReIEBd-Q	1YCeqlDI0ggsbByH3RRhw	Di Fara Pizza	Las Vegas	Restaurants, Italian	Mike 9538

- 'Oldest (yelping since)' influencer

user_id	business_id	name	city	categories	name yelping_since
c6HT44PKCaKqzN_BdgKPCw	u8C8pRvaHkg3PgDrsUHQ	Papa Del's Pizza	Champaign	Food Delivery Ser...	Russel 2004-10-12 08:40:43

- Influencer with the most 'review counts'

user_id	business_id	name	city	categories	name review_count
8k3a0-mPeyhbRSHUuca5aA	607-wkCpc1KF75JZLOTcMw	Circus Circus	Las Vegas	Arts & Entertainment, Restaurants	Victor 13278
RtGqdBvBcJcu5dUqwfzA	oUX2bYbajqST-urKb0HG6w	Loftti Cafe	Las Vegas	Desserts, Juice Bars & Smoothies...	Shila 12390

- Most 'Useful' influencer in Las Vegas

user_id	business_id	name	city	categories	name useful
--2vR0DIsmQ6Wfc5zKWigw	uanC140GcImHLG1_AT4JHq	Treasure Island	Las Vegas	Hair Salons, Arts...	Harald 154202

## III. Cities reviewed by the strongest influencers

- Las Vegas on top:

user_id	business_id	name	city	name review_count
8k3a0-mPeyhbRSHUuca5aA	607-wkCpc1KF75JZLOTcMw	Circus Circus Las Vegas Hotel and Casino	Las Vegas	Victor 13278
RtGqdBvBcJcu5dUqwfzA	oUX2bYbajqST-urKb0HG6w	Loftti Cafe	Las Vegas	Shila 12390

# Most Popular Influencers

- Las Vegas has the Influencer with the most 'Useful' reviews (business review count)

user_id	business_id	name	city	review_count	name useful
2vr80Ism06WfcSzKWigw	I88zLlGra0g9LU7qQVLpyg	Fashion Show	Las Vegas	739	Harald 154202
2vr80Ism06WfcSzKWigw	uanCi40Gc1mHLG_AT4JhQ	Treasure Island	Las Vegas	2487	Harald 154202
2vr80Ism06WfcSzKWigw	7dHYudt600IjiaxkSvv3lQ	In-N-Out Burger	Las Vegas	417	Harald 154202

- Champaign has the 'oldest (yelping since)' influencer (business review count)

user_id	business_id	name	city	review_count	name	yelping_since
c6HT44PKCaXqzN_BdgKPCw	u8CBpRvaHXg3PgDrsUJHJQ	Papa Del's Pizza	Champaign	402	Russel	2004-10-12 08:40:43

- Las Vegas has the influencer with the most 'fans' (business review count)

user_id	business_id	name	city	review_count	name fans
37cpUoM8hLkSQfReIEBd-Q	0qet57CmMA5qUm6gPFUTpg	Di Fara Pizza	Las Vegas	150	Mike 9538

- The most 'fans'

user_id	name	fans
37cpUoM8hLkSQfReIEBd-Q	Mike	9538
hizGc5W1tBHPghM5YKCAtg	Katie	2964
eKUGKQRE-Ywi5dY55_zChg	Cherylynn	2434
iLjMdZi0Tm7DQxX1C1_2dg	Ruggy	2383
j14WgRoU_-2ZE1aw1dXrJg	Daniel	2132

- The oldest 'yelping since'

user_id	name	yelping_since
c6HT44PKCaXqzN_BdgKPCw	Russel	2004-10-12 08:40:43
nkN_do3fJ9xekchVC-v68A	Jeremy	2004-10-12 08:46:43
wqoXYLWmpKEH0YvTmHBsJQ	Michael	2004-10-12 08:51:07
sE3ge33huDcNJGW3V4obww	Ken	2004-10-12 09:16:01
5i0Hz6pHmXi9SoB5qomRWQ	Nader	2004-10-12 17:42:24

- Most recent Elite influencer

user_id	business_id	name	city	name	elite
3Fmj7MfGfsUUK1kTWCSL_g	D5oLn4j7eezCAo0suYr8jA	ND Sushi & Grill	Toronto	Matthew	2018

## V. Top influencers:

- The most 'Useful'

user_id	name	useful
2vr80Ism06WfcSzKWigw	Harald	154202
JjXuiru1_ONzDkYVrHn0aw	Richard	99162
W7DHyQlY_kXls2iXt-_2Ag	Maggie	89792
Hi10sGS2NxQH3NlyWSZ1oA	Fox	89418
ax75nXOTIpatbsmqHLqVow	Rohlin	81003

- The most 'review counts'

user_id	name	review_count	average_stars
8k3a0-mPeyhbr5HUucA5aA	Victor	13278	3.28

# Recommendation system


We want to recommend restaurants to users

- 2 main type of recommenders:
  - Content-based recommendations (good to recommend when user has no history)
  - Collaborative filtering (good when user has a history)
- Yelp Dataset:
  - Restaurant data contains very limited information
  - Very few reviews compared to the number of restaurants and user(sparse matrix)
- Chosen approach:
  - we use collaborative filtering through matrix factorization to predict user ratings based on past ratings
  - Use ALS in SparkML

## ALS algorithm:

- Factorize a ratings matrix  $R$  into two latent factor matrices which when multiplied back, will give a approximation of the original ratings matrix. In the approximation matrix, all the cells will be filled by an estimated rating.
- if we want to predict how user  $K$  might rate product  $J$  we just multiply those two vector together
- Cost function:

$$\min_{x_*, y_*} \sum_{r_{u,i} \text{ is known}} (r_{ui} - x_u^T y_i)^2 + \lambda(\|x_u\|^2 + \|y_i\|^2)$$




- It's alternating because the process that generates those matrices  $U$  and  $P$  is done first by fixing  $U$  optimizing for  $P$  and then fixing  $P$  and optimizing for  $U$  and we repeat that process alternately. → effective performance
- Using ASL-WR to avoid overfitting

# Cleaning up the data even more

Reduce data to reduce memory consumption and speed up the process:

- Consider Toronto since it has the most number of restaurants
- Only keep restaurants with over 100 reviews
- Only keep users with over 10 reviews
- 6900 restaurant ---> 895 restaurants
- 80000 users ---> 5458 users

Cumulative number of restaurants per number of reviews



Cumulative number of users per number of reviews





# Tuning the recommender



ALS relies on 3 hyperparameters:

- Number of latent features
- Number of iterations
- Lambda of regularization

The best values of the hyperparameters are the one which minimize the RSE cost function.

We run 420 experiments with different hyperparameters in 5 hours.

The best set of parameter is ....[show map]

# Experiment result


Running in.... Mins

RMSE = ....

Sample of recommendation

# Demo

- Front-end:
  - React Framework: Popular Javascript Framework
  - Styling: Materialize
  - Map Library: Leaflet
  - Chart Library: VIS
- Backend:
  - Flask: Popular Python Framework for Web Services



# Contribution

Amaan: preprocessing data using Spark SQL

TJ: aggregation data using Spark SQL

Kathia: aggregation data using Spark Core (low level API)

Truc: Recommendation with Spark ML

All: Web application

# Future works

- Try with larger data using a large cluster of machines
- Show recommendation to real users and evaluate performance using A/B testing
- Try other models: hybrid, neural network
- Use database as backend

Questions?

